

## Briefly introduction to the algorithm or overall system

In this tutorial, we are going to learn how to use Vitis HLS to build, analyze, and optimize a hardware kernel. There are a few ways to do it, such as pipelining, using dual-port RAMs, array partitioning, and dataflow optimization.

The algorithm implemented in hardware is discrete cosine transform, DCT. It can process data into frequency domain, similar to Fourier transform. What makes them different is that DCT only uses real numbers. DCT is often used to process signals and images. Assume that we have a matrix with size  $N \times N$ , and the values are denoted by  $f(x, y)$ . The formula is shown below.

$$D(i, j) = \frac{1}{2\sqrt{2N}} C(i)C(j) \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) \cos \frac{(2x+1)i\pi}{2N} \cos \frac{(2y+1)j\pi}{2N}$$
$$C(i) = \begin{cases} \frac{1}{\sqrt{2}}, & \text{if } i = 0 \\ 1, & \text{otherwise} \end{cases}$$

## Explain the original code/system/pragmas and how you implement it

```
void dct(short input[N], short output[N]) {  
  
    short buf_2d_in[DCT_SIZE][DCT_SIZE];  
    short buf_2d_out[DCT_SIZE][DCT_SIZE];  
  
    // Read input data. Fill the internal buffer.  
    read_data(input, buf_2d_in);  
  
    dct_2d(buf_2d_in, buf_2d_out);  
  
    // Write out the results.  
    write_data(buf_2d_out, output);  
    //write_data(buf_2d_in, output);  
}
```

This is the top function. It is obvious that the kernel will read data, process the data, and then write the data. Note that DCT\_SIZE is specified as 8 in the design.

```
void read_data(short input[N], short buf[DCT_SIZE][DCT_SIZE])  
{  
    int r, c;  
  
    RD_Loop_Row:  
    for (r = 0; r < DCT_SIZE; r++) {  
        RD_Loop_Col:  
        for (c = 0; c < DCT_SIZE; c++)  
            buf[r][c] = input[r * DCT_SIZE + c];  
    }  
}
```

This is the function that reads data. We can see that it reads a 1d array into a 2d

array.

```
void write_data(short buf[DCT_SIZE][DCT_SIZE], short output[N])
{
    int r, c;

WR_Loop_Row:
    for (r = 0; r < DCT_SIZE; r++) {
WR_Loop_Col:
        for (c = 0; c < DCT_SIZE; c++)
            output[r * DCT_SIZE + c] = buf[r][c];
    }
}
```

This is the function that writes data. We can see that it writes a 2d array into a 1d array.

```
void dct_2d(dct_data_t in_block[DCT_SIZE][DCT_SIZE],
           dct_data_t out_block[DCT_SIZE][DCT_SIZE])
{
    dct_data_t row_outbuf[DCT_SIZE][DCT_SIZE];
    dct_data_t col_outbuf[DCT_SIZE][DCT_SIZE], col_inbuf[DCT_SIZE][DCT_SIZE];
    unsigned i, j;

    // DCT rows
WR_DCT_Loop:
    for(i = 0; i < DCT_SIZE; i++) {
        dct_1d(in_block[i], row_outbuf[i]);
    }
    // Transpose data in order to re-use 1D DCT code
Xpose_Row_Outer_Loop:
    for (j = 0; j < DCT_SIZE; j++)
Xpose_Row_Inner_Loop:
        for(i = 0; i < DCT_SIZE; i++)
            col_inbuf[j][i] = row_outbuf[i][j];
    // DCT columns
Col_DCT_Loop:
    for (i = 0; i < DCT_SIZE; i++) {
        dct_1d(col_inbuf[i], col_outbuf[i]);
    }
    // Transpose data back into natural order
Xpose_Col_Outer_Loop:
    for (j = 0; j < DCT_SIZE; j++)
Xpose_Col_Inner_Loop:
        for(i = 0; i < DCT_SIZE; i++)
            out_block[j][i] = col_outbuf[i][j];
}
```

This is the core function of DCT algorithm. The function is further divided into 4 parts, processing row, transposing, processing column, and transposing again. In this way, they can reuse the same function. Here, I will show how they divide the formula into row and column steps with  $N = 8$ .

$$D(i, j) = \frac{1}{2\sqrt{2N}} C(i)C(j) \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) \cos \frac{(2x+1)i\pi}{2N} \cos \frac{(2y+1)j\pi}{2N}$$

$$= \frac{1}{8} C(i)C(j) \sum_{x=0}^7 \sum_{y=0}^7 f(x, y) \cos \frac{(2x+1)i\pi}{16} \cos \frac{(2y+1)j\pi}{16}$$

$$= \frac{1}{2\sqrt{2}} C(i) \sum_{x=0}^7 \left[ \frac{1}{2\sqrt{2}} C(j) \sum_{y=0}^7 f(x, y) \cos \frac{(2y+1)j\pi}{16} \right] \cos \frac{(2x+1)i\pi}{16}$$

$$Temp(x, j) = \frac{1}{2\sqrt{2}} C(j) \sum_{y=0}^7 f(x, y) \cos \frac{(2y+1)j\pi}{16}$$

$$D(i, j) = \frac{1}{2\sqrt{2}} C(i) \sum_{x=0}^7 Temp(x, j) \cos \frac{(2x+1)i\pi}{16}$$

We can see that the formula has the similar operation for sum over x and sum over y. They declare another function to do the operation row by row, and transpose is required.

```
void dct_1d(dct_data_t src[DCT_SIZE], dct_data_t dst[DCT_SIZE])
{
    unsigned int k, n;
    int tmp;
    const dct_data_t dct_coeff_table[DCT_SIZE][DCT_SIZE] = {
#include "dct_coeff_table.txt"
    };

    DCT_Outer_Loop:
    for (k = 0; k < DCT_SIZE; k++) {
    DCT_Inner_Loop:
        for(n = 0, tmp = 0; n < DCT_SIZE; n++) {
            int coeff = (int)dct_coeff_table[k][n];
            tmp += src[n] * coeff;
        }
        dst[k] = DESCALE(tmp, CONST_BITS);
    }
}
```

In this function, they input a row, and calculate the accumulated sum (inner loop) for each element (outer loop). Here, a  $8 \times 8$  coefficient table, `dct_coeff_table.txt`, is used.

```
8192, 8192, 8192, 8192, 8192, 8192, 8192, 8192,
11363, 9633, 6436, 2260, -2260, -6436, -9632, -11362,
10703, 4433, -4433, -10703, -10703, -4433, 4433, 10703,
9633, -2260, -11362, -6436, 6436, 11363, 2260, -9632,
8192, -8192, -8192, 8192, 8192, -8191, -8191, 8192,
6436, -11362, 2260, 9633, -9632, -2260, 11363, -6436,
4433, -10703, 10703, -4433, -4433, 10703, -10703, 4433,
2260, -6436, 9633, -11362, 11363, -9632, 6436, -2260
```

In general, people convert floating numbers into integers before computing in hardware. Here, they implement the conversion by multiplying the floating numbers by  $2^{15}$ . Therefore, the coefficients can be calculated in this way.

$$Coef(j, y) = \frac{1}{2\sqrt{2}} C(j) \cos \frac{(2y+1)j\pi}{16} \times 2^{15}$$

## Examine the synthesis log to list what steps the tool takes during synthesis.

```
INFO: [SYN 201-201] Setting up clock 'default' with a period of 10ns.
INFO: [HLS 200-1611] Setting target device to 'xcu50-fsvh2104-2L-e'
INFO: [HLS 200-1505] Using flow_target 'vitis'
INFO: [HLS 200-1464] Running solution command: config_interface -m_axi_latency=64
INFO: [HLS 200-1464] Running solution command: config_interface -m_axi_alignment_byte_size=64
INFO: [HLS 200-1464] Running solution command: config_interface -m_axi_max_widen_bitwidth=512
INFO: [HLS 200-1464] Running solution command: config_interface -default_interface=kernel
INFO: [HLS 200-1464] Running solution command: config_rtl -register_reset_num=3
INFO: [HLS 200-1510] Running: set_part xcu50-fsvh2104-2L-e
INFO: [HLS 200-1510] Running: create_clock -period 10 -name default
INFO: [HLS 200-1510] Running: config_interface -m_axi_alignment_byte_size 64 -m_axi_latency 64 -m_axi_max_widen_bitwidth 512
INFO: [HLS 200-1510] Running: config_rtl -register_reset_num 3
INFO: [HLS 200-1510] Running: csynth_design
INFO: [HLS 200-111] Finished File checks and directory preparation: CPU user time: 0.01 seconds. CPU system time: 0.01 seconds. Elapsed time: 0.01 seconds; current allocated memory: 1.212 GB.
```

Project and solution initialization loads source and constraints files, and configures the active solution for synthesis.

```
INFO: [HLS 200-10] Analyzing design file 'Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp' ...
INFO: [HLS 200-111] Finished Source Code Analysis and Preprocessing: CPU user time: 2.01 seconds. CPU system time: 0.22 seconds. Elapsed time: 1.48 seconds; current allocated memory: 207.621 MB.
```

Start compilation reads source files into memory.

```
INFO: [HLS 200-777] Using interface defaults for 'Vitis' flow target.
```

Interface detection and setup reviews and generates port and block interfaces for the function.

```
INFO: [HLS 214-115] Multiple burst reads of length 2 and bit width 512 has been inferred on bundle 'gmem'. These burst requests might be further partitioned into multiple requests during RTL generation, based on max_read_burst_length or max_write_burst_length settings. (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:74:4)
INFO: [HLS 214-115] Multiple burst writes of length 2 and bit width 512 has been inferred on bundle 'gmem'. These burst requests might be further partitioned into multiple requests during RTL generation, based on max_read_burst_length or max_write_burst_length settings. (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:87:4)
```

Burst read and write analysis for ports/interfaces.

```
INFO: [HLS 200-10] Starting code transformations ...
INFO: [HLS 200-111] Finished Standard Transforms: CPU user time: 0.02 seconds. CPU system time: 0 seconds. Elapsed time: 0.07 seconds; current allocated memory: 209.516 MB.
```

Compiler transforms code to operations.

```
INFO: [HLS 200-10] Checking synthesizability ...
INFO: [HLS 200-111] Finished Checking Synthesizability: CPU user time: 0.03 seconds. CPU system time: 0 seconds. Elapsed time: 0.03 seconds; current allocated memory: 210.582 MB.
```

Performs Synthesizeability checks.

```
INFO: [XFORM 203-510] Pipelining loop 'RD_Loop_Row' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:71) in function 'dct' automatically.
INFO: [XFORM 203-510] Pipelining loop 'DCT_Outer_Loop' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:21) in function 'dct' automatically.
INFO: [XFORM 203-510] Pipelining loop 'Xpose_Row_Inner_Loop' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:43) in function 'dct' automatically.
INFO: [XFORM 203-510] Pipelining loop 'DCT_Outer_Loop' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:21) in function 'dct' automatically.
INFO: [XFORM 203-510] Pipelining loop 'Xpose_Col_Inner_Loop' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:43) in function 'dct' automatically.
INFO: [XFORM 203-510] Pipelining loop 'WR_Loop_Row' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:84) in function 'dct' automatically.
```

Automatic pipelining of loops at tripcount threshold.

```
INFO: [XFORM 203-502] Unrolling all sub-loops inside loop 'RD_Loop_Row' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:71) in function 'dct' for pipelining.
INFO: [XFORM 203-502] Unrolling all sub-loops inside loop 'DCT_Outer_Loop' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:21) in function 'dct' for pipelining.
INFO: [XFORM 203-502] Unrolling all sub-loops inside loop 'DCT_Outer_Loop' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:21) in function 'dct' for pipelining.
INFO: [XFORM 203-502] Unrolling all sub-loops inside loop 'WR_Loop_Row' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:84) in function 'dct' for pipelining.
INFO: [HLS 200-489] Unrolling loop 'RD_Loop_Col' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:71) in function 'dct' completely with a factor of 8.
INFO: [HLS 200-489] Unrolling loop 'DCT_Inner_Loop' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:21) in function 'dct' completely with a factor of 8.
INFO: [HLS 200-489] Unrolling loop 'WR_Loop_Col' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:84) in function 'dct' completely with a factor of 8.
```

Unrolling loops, both automatic and user-directed.

```
INFO: [XFORM 203-102] Partitioning array 'buf_2d_in' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:97) in dimension 2 automatically.
INFO: [XFORM 203-102] Partitioning array 'dct_coeff_table' in dimension 2 automatically.
```

Array partitioning.

```
INFO: [XFORM 203-111] Balancing expressions in function 'dct' (Vitis-Tutorials/Getting Started/Vitis_HLS/reference-files/src/dct.cpp:94)...16 expression(s) balanced.
```

Balance expressions using associative and commutative properties.

```
INFO: [XFORM 203-541] Flattening a loop nest 'Row_DCT_Loop' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:43:13) in function 'dct'.
INFO: [XFORM 203-541] Flattening a loop nest 'Xpose_Row_Outer_Loop' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:43:16) in function 'dct'.
INFO: [XFORM 203-541] Flattening a loop nest 'Col_DCT_Loop' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:43:13) in function 'dct'.
INFO: [XFORM 203-541] Flattening a loop nest 'Xpose_Col_Outer_Loop' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:43:16) in function 'dct'.
```

Loop flattening to reduce loop hierarchy.

```
INFO: [HLS 200-472] Inferring partial write operation for 'buf_2d_in[0]' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:78:11)
INFO: [HLS 200-472] Inferring partial write operation for 'row_outbuf' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:34:14)
INFO: [HLS 200-472] Inferring partial write operation for 'col_inbuf' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:55:26)
INFO: [HLS 200-472] Inferring partial write operation for 'buf_2d_out' (Vitis-Tutorials/Getting_Started/Vitis_HLS/reference-files/src/dct.cpp:66:26)
```

Partial write detection (writing part of a memory word)

```
INFO: [HLS 200-10] Starting hardware synthesis ...
INFO: [HLS 200-10] Synthesizing 'dct' ...
INFO: [HLS 200-10] -----
INFO: [HLS 200-42] -- Implementing module 'dct_Pipeline_RD_Loop_Row'
INFO: [HLS 200-10] -----
INFO: [HLS 200-10] Starting scheduling ...
INFO: [SCHED 204-61] Pipelining loop 'RD_Loop_Row'.
INFO: [HLS 200-1470] Pipelining result : Target II = NA, Final II = 1, Depth = 3, loop 'RD_Loop_Row'
INFO: [SCHED 204-11] Finished scheduling.
INFO: [HLS 200-111] Finished Scheduling: CPU user time: 0.04 seconds. CPU system time: 0.01 seconds. Elapsed time: 0.1 seconds; current allocated memory: 286.422 MB.
```

Finish architecture synthesis, start scheduling.

```
INFO: [HLS 200-10] -----
INFO: [HLS 200-10] -- Generating RTL for module 'dct_Pipeline_Row_DCT_Loop_DCT_Outer_Loop'
INFO: [HLS 200-10] -----
INFO: [HLS 200-1030] Apply Unified Pipeline Control on module 'dct_Pipeline_Row_DCT_Loop_DCT_Outer_Loop' pipeline 'Row_DCT_Loop_DCT_Outer_Loop' pipeline type 'Loop pipeline'
INFO: [RTGEN 206-100] Generating core module 'mac_muladd_16s_14ns_29s_29_4_1': 1 instance(s).
INFO: [RTGEN 206-100] Generating core module 'mac_muladd_16s_15s_13ns_29_4_1': 1 instance(s).
INFO: [RTGEN 206-100] Generating core module 'mac_muladd_16s_15s_29ns_29_4_1': 1 instance(s).
INFO: [RTGEN 206-100] Generating core module 'mac_muladd_16s_15s_29s_29_4_1': 2 instance(s).
INFO: [RTGEN 206-100] Generating core module 'mul_mul_16s_15s_29_4_1': 3 instance(s).
INFO: [RTGEN 206-100] Finished creating RTL model for 'dct_Pipeline_Row_DCT_Loop_DCT_Outer_Loop'.
INFO: [HLS 200-111] Finished Creating RTL model: CPU user time: 0.08 seconds. CPU system time: 0.02 seconds. Elapsed time: 0.18 seconds; current allocated memory: 293.012 MB.
```

End scheduling, generate RTL code.

Analyze the timing/performance/utilization

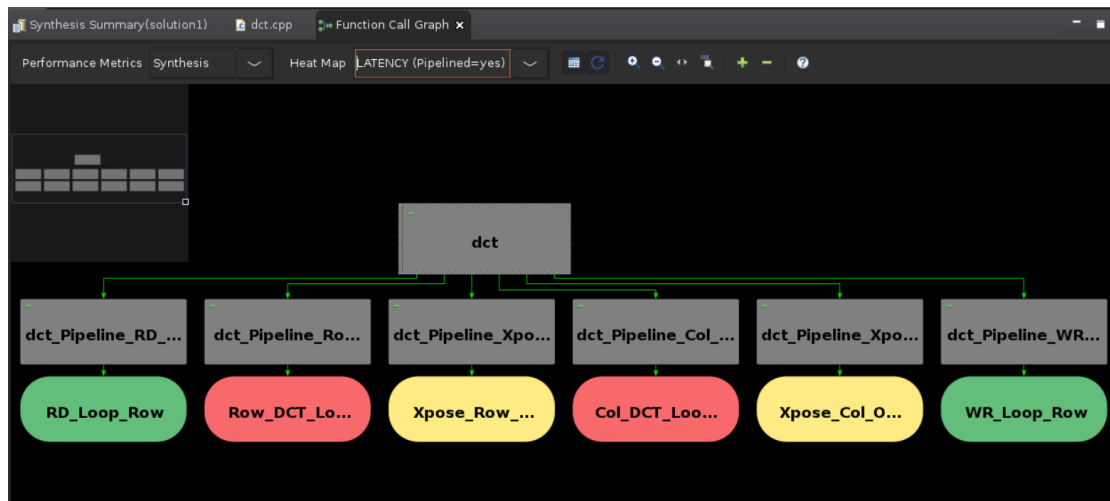
Modules & Loops	Issue Type	Slack	Latency (cycles)	Latency (ns)	Iteration Latency	Interval	Trip Count	Pipelined
+ dct	-	0.00	444	4.440e+03	-	445	-	no
+ dct_Pipeline_RD_Loop_Row	-	0.00	11	110.000	-	11	-	no
o RD_Loop_Row	-	7.30	9	90.000	3	1	8	yes
+ dct_Pipeline_Row_DCT_Loop_DCT_Outer_Loop	-	4.39	70	700.000	-	70	-	no
o Row_DCT_Loop_DCT_Outer_Loop	-	7.30	68	680.000	6	1	64	yes
+ dct_Pipeline_Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop	-	4.85	66	660.000	-	66	-	no
o Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop	-	7.30	64	640.000	2	1	64	yes
+ dct_Pipeline_Col_DCT_Loop_DCT_Outer_Loop	-	4.39	70	700.000	-	70	-	no
o Col_DCT_Loop_DCT_Outer_Loop	-	7.30	68	680.000	6	1	64	yes
+ dct_Pipeline_Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop	-	4.85	66	660.000	-	66	-	no
o Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop	-	7.30	64	640.000	2	1	64	yes
+ dct_Pipeline_WR_Loop_Row	-	0.00	11	110.000	-	11	-	no
o WR_Loop_Row	-	7.30	9	90.000	3	1	8	yes

This is the latency/loop table resulting from the original setting in the tutorial. To see how it calculate the latency of loops, I disable the pipeline of RD\_Loop\_Row. The result is shown below.

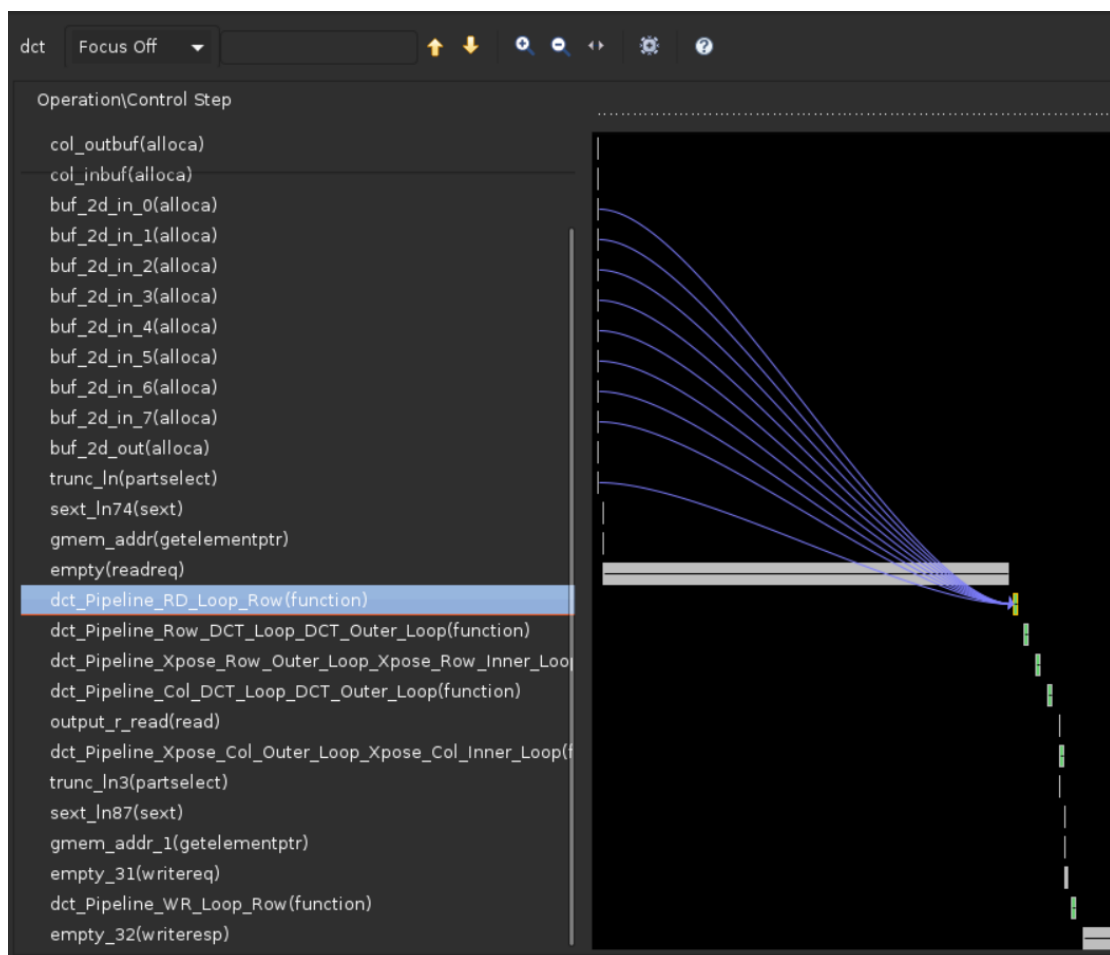
Modules & Loops	Issue Type	Slack	Latency (cycles)	Latency (ns)	Iteration Latency	Interval	Trip Count	Pipelined
o WR_Loop_Row	-	7.30	9	90.000	3	1	8	yes
o RD_Loop_Row	-	7.30	104	1.040e+03	13	-	8	no

We can see that if a loop is pipelined, the latency is small. The value is approximately  $Iteration\ Latency + Interval \times (Trip\ Count - 1) = 10$ . If a loop is not pipelined, the latency is large. The value is  $Iteration\ Latency \times Trip\ Count = 104$ .

From the latency table, we can also see that the critical blocks to optimize are the dct\_Pipeline\_Row\_DCT\_Loop\_DCT\_Outer\_Loop and dct\_Pipeline\_Col\_DCT\_Loop\_DCT\_Outer\_Loop, since their latency are largest. This can also be seen on Function Call Graph Viewer by setting heat map.

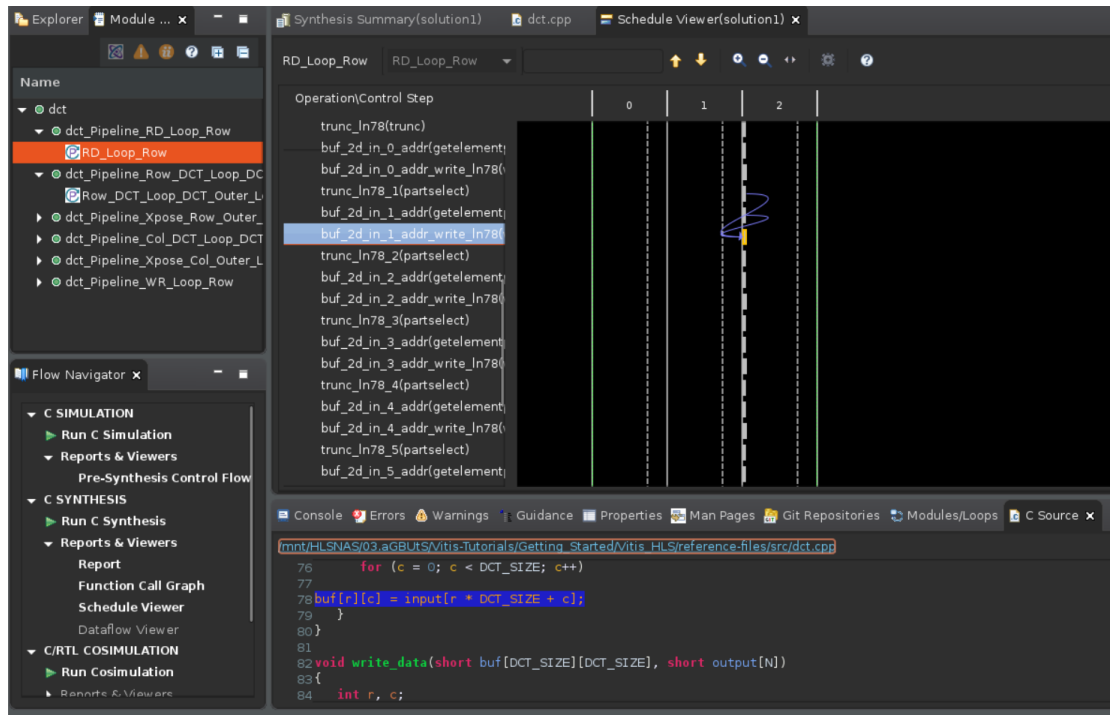


## Schedule viewer

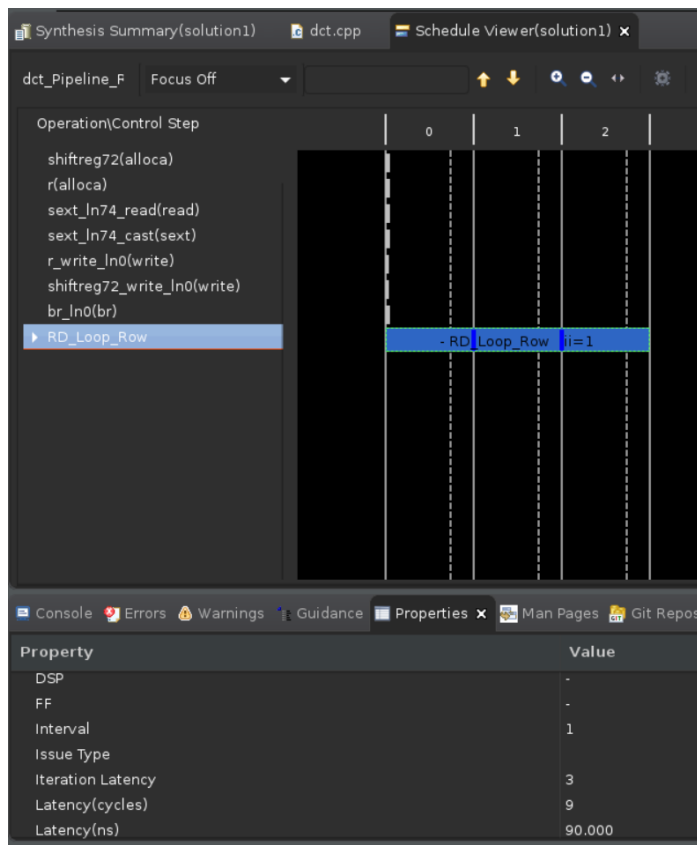


From the viewer of the top module, we can find those modules introduced above, and they are arranged in order. We can also see some blue lines representing the data dependencies.





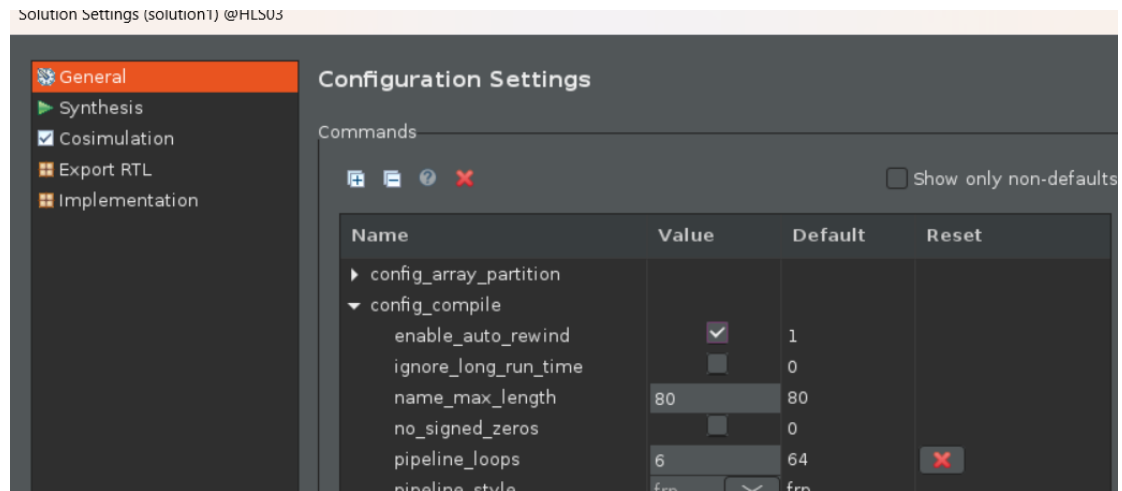
From the module hierarchy, we can select the modules and loops to view their schedule. Besides, we can also right click on the operation to see where it was called.



From the properties, we can also see the latency of selected operation.

If possible, share how you optimize the design and the trade-off you make

## 1. Configure the pipeline loops threshold




In this method, we are asked to modify pipeline\_loops as 6. It enables loops with an iteration count below 6 to be pipelined automatically.

Modules & Loops	Issue Type	Slack	Latency (cycles)	Latency (ns)	Iteration Latency	Interval	Trip Count	Pipelined
+ dct	-	0.00	2548	2.548e+04	-	2549	-	no
+ dct_Pipeline_Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop	-	4.85	66	660.000	-	66	-	no
o Xpose_Row_Outer_Loop_Xpose_Row_Inner_Loop	-	7.30	64	640.000	2	1	64	yes
+ dct_Pipeline_Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop	-	4.85	66	660.000	-	66	-	no
o Xpose_Col_Outer_Loop_Xpose_Col_Inner_Loop	-	7.30	64	640.000	2	1	64	yes
o RD_Loop_Row	-	7.30	120	1.200e+03	15	-	8	no
+ dct_Pipeline_RD_Loop_Col	-	5.91	10	100.000	-	10	-	no
o RD_Loop_Col	-	7.30	8	80.000	1	1	8	yes
o Row_DCT_Loop_DCT_Outer_Loop	-	7.30	1024	1.024e+04	16	-	64	no
+ dct_Pipeline_DCT_Inner_Loop	-	5.62	13	130.000	-	13	-	no
o DCT_Inner_Loop	-	7.30	11	110.000	5	1	8	yes
o Col_DCT_Loop_DCT_Outer_Loop	-	7.30	1024	1.024e+04	16	-	64	no
+ dct_Pipeline_DCT_Inner_Loop1	-	5.62	13	130.000	-	13	-	no
o DCT_Inner_Loop1	-	7.30	11	110.000	5	1	8	yes
o WR_Loop_Row	-	7.30	104	1.040e+03	13	-	8	no
+ dct_Pipeline_WR_Loop_Col	-	5.91	10	100.000	-	10	-	no
o WR_Loop_Col	-	7.30	8	80.000	2	1	8	yes

After synthesis, we can see that the latency increases a lot, and that many loops are not pipelined, resulting to large latency. By the way, increasing the value does not help in this case, since there is no loop with iteration counter larger than the default value, 64.

## 2. Configure the pipeline initiation interval

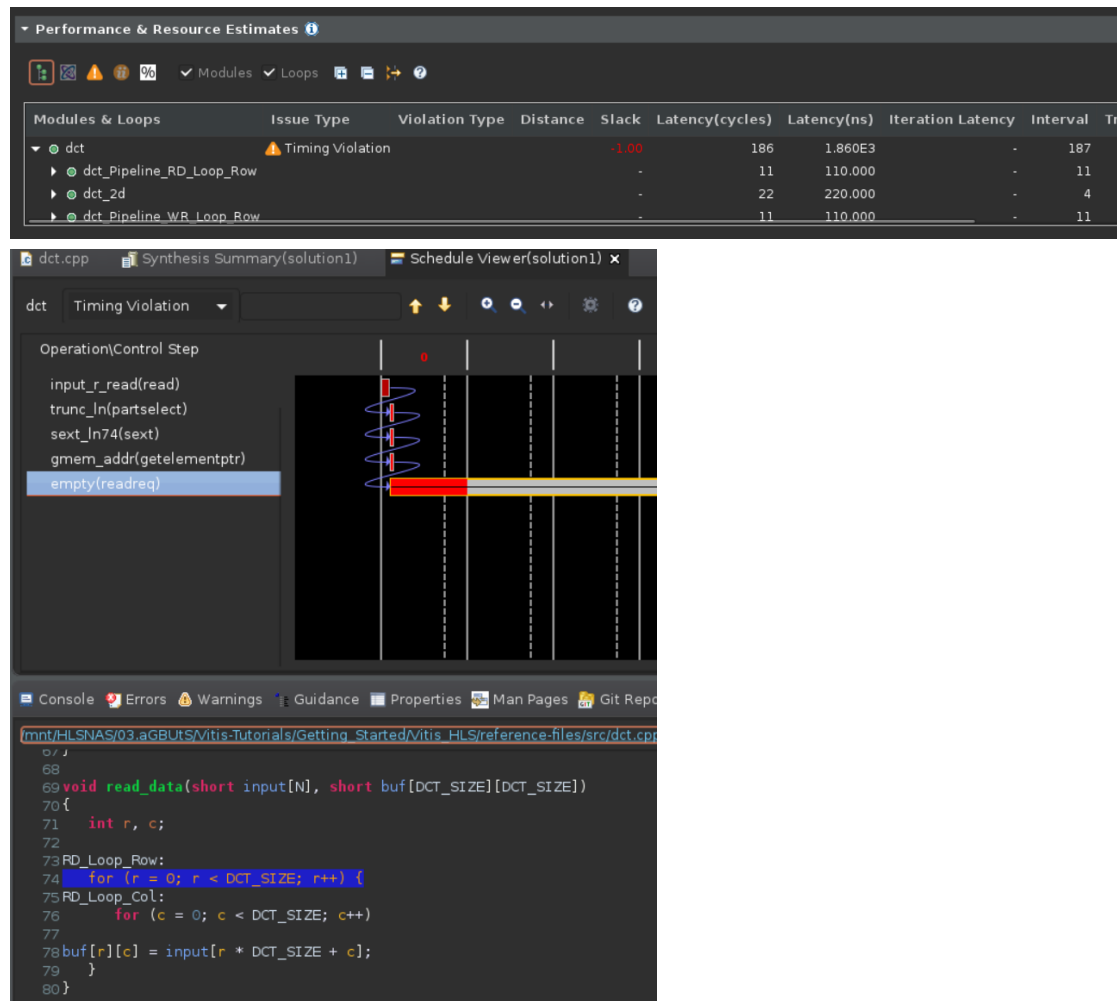
In the tutorial, we are asked to modify the II of dct\_2d as 4.




After synthesis, although the overall latency is decreased, there is timing



violation.



We can see that the timing violation occurred during reading data. I tried to figure out how to fix it by adding some directives, but end up failing. I tried to modify the II as 2, and it works well. The figure below is the result performance. We can see that the latency is greatly decreased.



### 3. Assign dual-port RAMs with BIND\_STORAGE

In the tutorial, we are asked to use BRAMs to store col\_inbuf and buf\_2d\_out. After synthesis, the result is shown in the figure below.

Performance & Resource Estimates							
Modules & Loops							
Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)		
				664	6.640E3		
				11	110.000		
				70	700.000		
				66	660.000		
	II Violation			263	2.630E3		
				66	660.000		
	II Violation			38	380.000		

We can see that the total latency increases a lot. Besides, II violation is issued (II violation does not influence the co-simulation result. In HLS, II violation will be issued if II is not 1.) Take a look at the report, and we can see that it uses more resource as well. Therefore, I think that this method is not for this application.

Original method	This optimization method
<pre> ===== == Utilization Estimates ===== * Summary: +-----+-----+-----+-----+-----+-----+   Name   BRAM_18K   DSP   FF   LUT   URAM   +-----+-----+-----+-----+-----+-----+   DSP   -   -   -   -   -     Expression   -   -   0   2   -     FIFO   -   -   -   -   -     Instance   0   16   6689   9847   0     Memory   0   -   329   340   0     Multiplexer   -   -   -   1748   -     Register   -   -   277   -   -   +-----+-----+-----+-----+-----+-----+   Total   0   16   7206   11937   0   +-----+-----+-----+-----+-----+-----+   Available SLR   1344   2976   871680   435840   320   +-----+-----+-----+-----+-----+-----+   Utilization SLR (%)   0   ~0   ~0   2   0   +-----+-----+-----+-----+-----+-----+   Available   2688   5952   1743360   871680   640   +-----+-----+-----+-----+-----+-----+   Utilization (%)   0   ~0   ~0   1   0   ===== </pre>	<pre> ===== == Utilization Estimates ===== * Summary: +-----+-----+-----+-----+-----+-----+   Name   BRAM_18K   DSP   FF   LUT   URAM   +-----+-----+-----+-----+-----+-----+   DSP   -   -   -   -   -     Expression   -   -   0   2   -     FIFO   -   -   -   -   -     Instance   0   16   6820   10110   0     Memory   2   -   288   306   0     Multiplexer   -   -   -   1640   -     Register   -   -   277   -   -   +-----+-----+-----+-----+-----+-----+   Total   2   16   7385   12058   0   +-----+-----+-----+-----+-----+-----+   Available SLR   1344   2976   871680   435840   320   +-----+-----+-----+-----+-----+-----+   Utilization SLR (%)   ~0   ~0   ~0   2   0   +-----+-----+-----+-----+-----+-----+   Available   2688   5952   1743360   871680   640   +-----+-----+-----+-----+-----+-----+   Utilization (%)   ~0   ~0   ~0   1   0   ===== </pre>

I also tried to remove the II violation. I think that violation happened since the bandwidth of BRAM is not large enough. Therefore, I try to partition the array (the next optimization method). The result is shown below.

Performance & Resource Estimates							
Modules & Loops							
Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)		
				475	4.750E3		
				11	110.000		
				70	700.000		
				66	660.000		
				70	700.000		
				66	660.000		
	II Violation			38	380.000		

We can see that the latency is reduced (664 -> 475), but still larger than that of the original method (444). To remove the II violation of writing output, I enlarge the II of WR\_Loop\_ROW as 2. The result is shown below.

Performance & Resource Estimates						
Modules & Loops						
Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	
-	-	-	-	459	4.590E3	
-	-	-	-	11	110.000	
-	-	-	-	70	700.000	
-	-	-	-	66	660.000	
-	-	-	-	70	700.000	
-	-	-	-	66	660.000	
-	-	-	-	18	180.000	

No II violation is issued, and the latency decreases a little (still larger than that of the original method).

Original method	This optimization method
<pre> ===== == Utilization Estimates ===== * Summary: +-----+-----+-----+-----+-----+-----+   Name   BRAM_18K   DSP   FF   LUT   URAM   +-----+-----+-----+-----+-----+-----+   DSP   -   -   -   -   -     Expression   -   -   0   2   -     FIFO   -   -   -   -   -     Instance   0   16   6689   9847   0     Memory   0   -   329   340   0     Multiplexer   -   -   -   1748   -     Register   -   -   277   -   -   +-----+-----+-----+-----+-----+-----+   Total   0   16   7206   11937   0   +-----+-----+-----+-----+-----+-----+   Available SLR   1344   2976   871680   435840   320   +-----+-----+-----+-----+-----+-----+   Utilization SLR (%)   0   ~0   ~0   2   0   +-----+-----+-----+-----+-----+-----+   Available   2688   5952   1743360   871680   640   +-----+-----+-----+-----+-----+-----+   Utilization (%)   0   ~0   ~0   1   0   +-----+-----+-----+-----+-----+-----+ </pre>	<pre> ===== == Utilization Estimates ===== * Summary: +-----+-----+-----+-----+-----+-----+   Name   BRAM_18K   DSP   FF   LUT   URAM   +-----+-----+-----+-----+-----+-----+   DSP   -   -   -   -   -     Expression   -   -   0   2   -     FIFO   -   -   -   -   -     Instance   0   16   6739   10378   0     Memory   16   -   288   306   0     Multiplexer   -   -   -   2871   -     Register   -   -   2333   -   -   +-----+-----+-----+-----+-----+-----+   Total   16   16   9360   13557   0   +-----+-----+-----+-----+-----+-----+   Available SLR   1344   2976   871680   435840   320   +-----+-----+-----+-----+-----+-----+   Utilization SLR (%)   1   ~0   1   3   0   +-----+-----+-----+-----+-----+-----+   Available   2688   5952   1743360   871680   640   +-----+-----+-----+-----+-----+-----+   Utilization (%)   ~0   ~0   ~0   1   0   +-----+-----+-----+-----+-----+-----+ </pre>

Again, compare their utilization, and we can see that the original method looks better.

#### 4. Assign an array partition

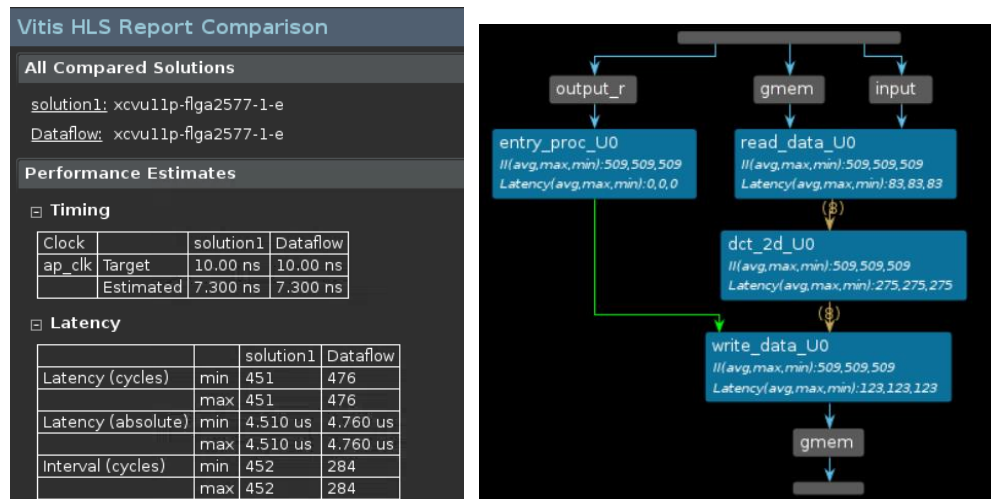
In the tutorial, we are asked to partition the arrays, col\_inbuf and buf\_2d\_out. The result is shown in the figures below.

Original method	Array partition
<pre> +-----+-----+-----+-----+-----+-----+-----+-----+   Modules &amp; Loops   Issue Type   Slack   Latency (cycles)   Latency (ns)   Iteration Latency   Interval   Trip Count   Pipelined   +-----+-----+-----+-----+-----+-----+-----+-----+   + dct   -   0.00   444   4.440E+03   -   445   -   no     + dct_Pipeline_RD_Loop_Row   -   0.00   11   110.000   -   11   -   no     o RD_Loop_Row   -   7.30   9   90.000   3   1   8   yes     + dct_Pipeline_Row_DCT_Loop_DCT_Out_Loop   -   4.39   70   700.000   -   70   -   no     o Row_DCT_Loop_DCT_Out_Loop   -   7.30   68   680.000   6   1   64   yes     + dct_Pipeline_Xpose_Row_Out_Loop_Xpose_Row_Inner_Loop   -   4.85   66   660.000   -   66   -   no     o Xpose_Row_Out_Loop_Xpose_Row_Inner_Loop   -   7.30   64   640.000   2   1   64   yes     + dct_Pipeline_Col_DCT_Loop_DCT_Out_Loop   -   4.39   70   700.000   -   70   -   no     o Col_DCT_Loop_DCT_Out_Loop   -   7.30   68   680.000   6   1   64   yes     + dct_Pipeline_Xpose_Col_Out_Loop_Xpose_Col_Inner_Loop   -   4.85   66   660.000   -   66   -   no     o Xpose_Col_Out_Loop_Xpose_Col_Inner_Loop   -   7.30   64   640.000   2   1   64   yes     + dct_Pipeline_WR_Loop_Row   -   0.00   11   110.000   -   11   -   no     o WR_Loop_Row   -   7.30   9   90.000   3   1   8   yes   +-----+-----+-----+-----+-----+-----+-----+-----+ </pre>	<pre> +-----+-----+-----+-----+-----+-----+-----+-----+   Modules &amp; Loops   Issue Type   Slack   Latency (cycles)   Latency (ns)   Iteration Latency   Interval   Trip Count   Pipelined   +-----+-----+-----+-----+-----+-----+-----+-----+   + dct   -   0.00   451   4.510E+03   -   452   -   no     + dct_Pipeline_RD_Loop_Row   -   0.00   11   110.000   -   11   -   no     o RD_Loop_Row   -   7.30   9   90.000   3   1   8   yes     + dct_Pipeline_Row_DCT_Loop_DCT_Out_Loop   -   4.39   70   700.000   -   70   -   no     o Row_DCT_Loop_DCT_Out_Loop   -   7.30   68   680.000   6   1   64   yes     + dct_Pipeline_Xpose_Row_Out_Loop_Xpose_Row_Inner_Loop   -   4.85   66   660.000   -   66   -   no     o Xpose_Row_Out_Loop_Xpose_Row_Inner_Loop   -   7.30   64   640.000   2   1   64   yes     + dct_Pipeline_Col_DCT_Loop_DCT_Out_Loop   -   4.39   70   700.000   -   70   -   no     o Col_DCT_Loop_DCT_Out_Loop   -   7.30   68   680.000   6   1   64   yes     + dct_Pipeline_Xpose_Col_Out_Loop_Xpose_Col_Inner_Loop   -   4.85   66   660.000   -   66   -   no     o Xpose_Col_Out_Loop_Xpose_Col_Inner_Loop   -   7.30   64   640.000   2   1   64   yes     + dct_Pipeline_WR_Loop_Row   -   0.00   10   100.000   -   10   -   no     o WR_Loop_Row   -   7.30   8   80.000   2   1   8   yes   +-----+-----+-----+-----+-----+-----+-----+-----+ </pre>

We can see that the iteration latency of WR\_Loop\_Row decreases. However, I have no idea why the overall latency increases.

## 5. Dataflow optimization

In the tutorial, we are asked to add pragma DATAFLOW to the directives. This pragma enables task-level pipelining. The HLS tool will analyze the dataflow between functions and see whether they can overlap or not. The comparison result is shown in the figure below.



We can see that the interval between functions decreases a lot.

After running co-simulation, we can see the dataflow graph. From the graph, we can see that the three blocks, dct\_2d, read\_data, write\_data are not running in parallel. From the process table below, we can see some information.

Process	Channel							
Name	Cosim Category	Cosim Stalling Time	FIFO EMPTY	FIFO FULL	Cosim Stall No Start	Cosim Stall No Continue	Cosim AVG II	
entry_proc_U0	none	0.00%	0.00%	0.00%	0.00%	0.00%	509	
read_data_U0	none	0.00%	0.00%	0.00%	0.00%	0.00%	509	
dct_2d_U0	none	0.00%	0.00%	0.00%	0.00%	0.00%	509	
write_data_U0	none	2.35%	0.00%	0.00%	0.00%	2.35%	509	

Cosim Max II	Cosim Min II	Cosim AVG Latency	Cosim Max Latency	Cosim Min Latency	Cosim Distribution Graph
509	509	0	0	0	<a href="#">Link</a>
509	509	83	83	83	<a href="#">Link</a>
509	509	275	275	275	<a href="#">Link</a>
509	509	123	123	123	<a href="#">Link</a>

From the channel table below, we can see the channel type and the bit width.

Process	Channel											
Name	Cosim Category	FIFO EMPTY	FIFO FULL	Cosim Max Depth	Depth	Type	Sub-Type	BitWidth	Producer	Consumer	Cosim Distribut	
buf_2d_in_0	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	
buf_2d_in_1	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	
buf_2d_in_2	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	
buf_2d_in_3	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	
buf_2d_in_4	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	
buf_2d_in_5	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	
buf_2d_in_6	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	
buf_2d_in_7	N/A	N/A	N/A	N/A	0	PIPO	PIPO	16	read_data_U0	dct_2d_U0	N/A	

Can it be further optimized? I tried the pipeline method taught in second method, and removed the array partitioning. The result is shown in the figure

below.

Performance Estimates			
▢ Timing			
Clock		solution1	DATAFLOW
ap_clk	Target	10.00 ns	10.00 ns
	Estimated	7.300 ns	7.300 ns
▢ Latency			
		solution1	DATAFLOW
Latency (cycles)	min	444	179
	max	444	179
Latency (absolute)	min	4.440 us	1.790 us
	max	4.440 us	1.790 us
Interval (cycles)	min	445	84
	max	445	84

We can see that the interval is now only 84.

### Explain what you observed and learned

In this lab, I learned some tools to help us analysis our design, such as schedule viewer, function call graph, and dataflow viewer. I also learned some optimization methods, such as pipelining, using dual-port RAMs, array partitioning, and dataflow optimization.

### Explain what problem you encountered and how you solved it

Timing violation in the second optimization method.

Il violation in the third optimization method.

Solutions have been introduced in the above section.

### Submit your work to Github

link: <https://github.com/JasonYanggg/HLS>